

Airline Business Intelligence – Phase 2 Data Quality Snapshot

This notebook is a quick **smoke test** for Phase 2 of the pipeline:

- OpenFlights reference data (`airports`, `airlines`)
- Synthetic flight schedule (`flights`)
- BTS on-time performance (`flight_performance`)
- Synthetic customers & loyalty (`passengers`, `loyalty_accounts`, `miles_transactions`)
- Synthetic revenue (`bookings`, `payments`)

Assumptions:

- PostgreSQL is running
- The `airline` schema has been migrated
- The environment variable `DATABASE_URL` points at the `airline_bi` database

```
In [1]: import os
from textwrap import dedent

import pandas as pd
from sqlalchemy import create_engine

db_url = os.getenv("DATABASE_URL")
if not db_url:
    raise RuntimeError("DATABASE_URL is not set - export it before running t
engine = create_engine(db_url, future=True)
```

```
In [2]: import os
from textwrap import dedent

import pandas as pd
from sqlalchemy import create_engine

db_url = os.getenv("DATABASE_URL")
if not db_url:
    raise RuntimeError("DATABASE_URL is not set - export it before running t
engine = create_engine(db_url, future=True)
```

1. Row counts by table

```
In [4]: def run_sql(query: str) -> pd.DataFrame:
    """Run a SQL query and return a pandas DataFrame."""

```

```

with engine.begin() as con:
    return pd.read_sql_query(dedent(query), con)

row_counts_sql = """
    SELECT 'airports' AS table_name, COUNT(*) AS row_count FROM airports
    UNION ALL
    SELECT 'airlines', COUNT(*) FROM airline.airlines
    UNION ALL
    SELECT 'flights', COUNT(*) FROM airline.flights
    UNION ALL
    SELECT 'flight_performance', COUNT(*) FROM airline.flight_performance
    UNION ALL
    SELECT 'passengers', COUNT(*) FROM airline.passengers
    UNION ALL
    SELECT 'loyalty_accounts', COUNT(*) FROM airline.loyalty_accounts
    UNION ALL
    SELECT 'miles_transactions', COUNT(*) FROM airline.miles_transactions
    UNION ALL
    SELECT 'bookings', COUNT(*) FROM airline.bookings
    UNION ALL
    SELECT 'payments', COUNT(*) FROM airline.payments
    ORDER BY table_name;
"""

row_counts = run_sql(row_counts_sql)
row_counts

```

Out [4]:

	table_name	row_count
0	airlines	5733
1	airports	7697
2	bookings	40000
3	flight_performance	22595
4	flights	5000
5	loyalty_accounts	3000
6	miles_transactions	10576
7	passengers	5000
8	payments	40000

2. Basic null checks on key columns

In [6]:

```

null_checks_sql = """
SELECT 'airports' AS table_name,
       COUNT(*) AS total_rows,
       SUM(CASE WHEN iata_code IS NULL THEN 1 ELSE 0 END) AS metric_1_nulls,
       SUM(CASE WHEN icao_code IS NULL THEN 1 ELSE 0 END) AS metric_2_nulls
FROM airline.airports
UNION ALL
"""

```

```

SELECT 'airlines' AS table_name,
       COUNT(*)   AS total_rows,
       SUM(CASE WHEN iata_code IS NULL THEN 1 ELSE 0 END) AS metric_1_nulls,
       SUM(CASE WHEN icao_code IS NULL THEN 1 ELSE 0 END) AS metric_2_nulls
  FROM airline.airlines
UNION ALL
SELECT 'passengers' AS table_name,
       COUNT(*)   AS total_rows,
       SUM(CASE WHEN email IS NULL THEN 1 ELSE 0 END)           AS metric_1_r
       SUM(CASE WHEN state_or_country IS NULL THEN 1 ELSE 0 END) AS metric_2
  FROM airline.passengers
UNION ALL
SELECT 'bookings' AS table_name,
       COUNT(*)   AS total_rows,
       SUM(CASE WHEN passenger_id IS NULL THEN 1 ELSE 0 END) AS metric_1_nul
       SUM(CASE WHEN flight_id     IS NULL THEN 1 ELSE 0 END) AS metric_2_nul
  FROM airline.bookings
UNION ALL
SELECT 'payments' AS table_name,
       COUNT(*)   AS total_rows,
       SUM(CASE WHEN booking_id IS NULL THEN 1 ELSE 0 END) AS metric_1_nulls
       SUM(CASE WHEN amount_usd IS NULL THEN 1 ELSE 0 END) AS metric_2_nulls
  FROM airline.payments
ORDER BY table_name;
"""

null_checks = run_sql(null_checks_sql)
null_checks

```

Out[6]:

	table_name	total_rows	metric_1_nulls	metric_2_nulls
0	airlines	5733	4619	172
1	airports	7697	1625	0
2	bookings	40000	0	0
3	passengers	5000	0	0
4	payments	40000	0	0

In [7]:

```

flights_nulls_sql = """
SELECT
       COUNT(*) AS total_rows,
       SUM(CASE WHEN airline_id          IS NULL THEN 1 ELSE 0 END) AS airline_
       SUM(CASE WHEN aircraft_id        IS NULL THEN 1 ELSE 0 END) AS aircraft_
       SUM(CASE WHEN route_id          IS NULL THEN 1 ELSE 0 END) AS route_id_
       SUM(CASE WHEN origin_airport_id  IS NULL THEN 1 ELSE 0 END) AS origin_i
       SUM(CASE WHEN destination_airport_id IS NULL THEN 1 ELSE 0 END) AS dest_
  FROM airline.flights;
"""

flights_nulls = run_sql(flights_nulls_sql)
flights_nulls

```

Out [7]:	total_rows	airline_id_nulls	aircraft_id_nulls	route_id_nulls	origin_id_nulls	dest_i
0	5000	0	5000	5000	0	

3. Primary-key and foreign-key sanity checks

```
In [8]: pk_fk_sql = """
-- PK uniqueness checks
SELECT 'airports_pk' AS check_name,
       COUNT(*) - COUNT(DISTINCT airport_id) AS issue_count
FROM airline.airports
UNION ALL
SELECT 'airlines_pk',
       COUNT(*) - COUNT(DISTINCT airline_id)
FROM airline.airlines
UNION ALL
SELECT 'flights_pk',
       COUNT(*) - COUNT(DISTINCT flight_id)
FROM airline.flights
UNION ALL
SELECT 'passengers_pk',
       COUNT(*) - COUNT(DISTINCT passenger_id)
FROM airline.passengers
UNION ALL
SELECT 'bookings_pk',
       COUNT(*) - COUNT(DISTINCT booking_id)
FROM airline.bookings
UNION ALL
SELECT 'payments_pk',
       COUNT(*) - COUNT(DISTINCT payment_id)
FROM airline.payments
UNION ALL

-- FK checks: flights -> airlines / airports
SELECT 'flight_airline_fk',
       COUNT(*)
FROM airline.flights f
LEFT JOIN airline.airlines al ON f.airline_id = al.airline_id
WHERE al.airline_id IS NULL

UNION ALL

SELECT 'flight_origin_airport_fk',
       COUNT(*)
FROM airline.flights f
LEFT JOIN airline.airports a ON f.origin_airport_id = a.airport_id
WHERE a.airport_id IS NULL

UNION ALL

SELECT 'flight_destination_airport_fk',
```

```
COUNT(*)
FROM airline.flights f
LEFT JOIN airline.airports a ON f.destination_airport_id = a.airport_id
WHERE a.airport_id IS NULL

UNION ALL

-- bookings / payments FKs
SELECT 'booking_flight_fk',
       COUNT(*)
FROM airline.bookings b
LEFT JOIN airline.flights f ON b.flight_id = f.flight_id
WHERE f.flight_id IS NULL

UNION ALL

SELECT 'booking_passenger_fk',
       COUNT(*)
FROM airline.bookings b
LEFT JOIN airline.passengers p ON b.passenger_id = p.passenger_id
WHERE p.passenger_id IS NULL

UNION ALL

SELECT 'payment_booking_fk',
       COUNT(*)
FROM airline.payments p
LEFT JOIN airline.bookings b ON p.booking_id = b.booking_id
WHERE b.booking_id IS NULL;
.....
pk_fk_issues = run_sql(pk_fk_sql)
pk_fk_issues
```

Out [8]:

	check_name	issue_count
0	passengers_pk	0
1	flights_pk	0
2	airlines_pk	0
3	airports_pk	0
4	flight_airline_fk	0
5	flight_destination_airport_fk	0
6	flight_origin_airport_fk	0
7	booking_passenger_fk	0
8	booking_flight_fk	0
9	payment_booking_fk	0
10	bookings_pk	0
11	payments_pk	0

4. Business sanity checks

In [9]:

```
bookings_per_passenger_sql = """
SELECT
    COUNT(*) AS total_bookings,
    COUNT(DISTINCT passenger_id) AS distinct_passengers,
    MIN(bookings_per_pax) AS min_bookings_per_passenger,
    PERCENTILE_CONT(0.5) WITHIN GROUP (ORDER BY bookings_per_pax)
        AS median_bookings_per_passenger,
    MAX(bookings_per_pax) AS max_bookings_per_passenger
FROM (
    SELECT passenger_id, COUNT(*) AS bookings_per_pax
    FROM airline.bookings
    GROUP BY passenger_id
) sub;
"""

bookings_per_passenger = run_sql(bookings_per_passenger_sql)
```

Out [9]:

	total_bookings	distinct_passengers	min_bookings_per_passenger	median_bookings_per_passenger
0	4995	4995	1	

In [10]:

```
payments_summary_sql = """
SELECT
    COUNT(*) AS num_payments,
    SUM(amount_usd) AS total_revenue_usd,
```

```

        AVG(amount_usd) AS avg_ticket_price_usd,
        MIN(amount_usd) AS min_ticket_price_usd,
        MAX(amount_usd) AS max_ticket_price_usd
    FROM airline.payments;
.....
payments_summary = run_sql(payments_summary_sql)
payments_summary

```

Out[10]:

	num_payments	total_revenue_usd	avg_ticket_price_usd	min_ticket_price_usd	max_ticket_price_usd
0	40000	4539794.86	113.494872	72.0	1000.0

In [15]:

```

import matplotlib.pyplot as plt

# Simple proof-of-pipeline chart: row counts per table
df = row_counts.sort_values("row_count")

plt.figure(figsize=(8,5))
plt.barh(df["table_name"], df["row_count"])
plt.title("Airline BI – Loaded Table Row Counts (Phase 2 Complete)")
plt.xlabel("Rows")
plt.tight_layout()

plt.savefig("/Users/gracepolito/Public/Airline Business Intelligence Database/row_counts.png")
plt.show()

```

