

.....

Backfill dimension / audit tables derived from existing flight data.

This script:

- 1) Populates airline.routes from distinct (airline_id, origin_airport_id, destination_airport_id)
- 2) Optionally computes distance_nm using airport latitude / longitude
- 3) Seeds airline.aircraft with a small synthetic global fleet
- 4) Assigns aircraft_id to flights
- 5) Creates synthetic airline.flight_changes records for a subset of flights

Schema assumptions (from the DB):

airline.flights	
flight_id	bigint PK
airline_id	bigint
route_id	bigint (nullable)
aircraft_id	bigint (nullable)
origin_airport_id	bigint
destination_airport_id	bigint
flight_number	text
flight_date	date
scheduled_departure_utc	timestamp
scheduled_arrival_utc	timestamp
actual_departure_utc	timestamp
actual_arrival_utc	timestamp
delay_minutes	integer
delay_cause	text
status	enum (flight_status)
airline.routes	
route_id	bigint PK
airline_id	bigint
origin_airport_id	bigint
destination_airport_id	bigint
distance_nm	integer (nullable)
airline.airports	
airport_id	bigint
iata_code	varchar
icao_code	varchar
name	text
city	text
country	text
latitude	numeric -- degrees
longitude	numeric -- degrees
timezone	text
airline.aircraft	
aircraft_id	bigint PK
manufacturer	text
model	text
seat_capacity	integer
tail_number	text (nullable)
airline.flight_changes	
change_id	bigint PK
flight_id	bigint FK -> flights
old_aircraft_id	bigint
new_aircraft_id	bigint
reason	text
changed_at	timestamp

.....

```

import os
from sqlalchemy import create_engine, text

def get_db_url() -> str:
    url = os.getenv("DATABASE_URL") or os.getenv("AIRLINE_DB_DSN")
    if not url:
        raise RuntimeError(
            "Set DATABASE_URL or AIRLINE_DB_DSN in your environment.\n"
            "Example: postgresql+psycopg2://postgres:gpcool@localhost:5432/airline_bi"
        )
    return url

ENGINE = create_engine(get_db_url(), future=True, pool_pre_ping=True)

# -----
# 1. ROUTES
# -----


def backfill_routes_from_flights(compute_distance: bool = False) -> None:
    """Insert missing routes based on distinct flights."""

    print("◆ Backfilling airline.routes from airline.flights ...")

    insert_sql = text(
        """
        INSERT INTO airline.routes (
            airline_id,
            origin_airport_id,
            destination_airport_id,
            distance_nm
        )
        SELECT DISTINCT
            f.airline_id,
            f.origin_airport_id,
            f.destination_airport_id,
            NULL::integer AS distance_nm
        FROM airline.flights f
        LEFT JOIN airline.routes r
            ON r.airline_id      = f.airline_id
            AND r.origin_airport_id = f.origin_airport_id
            AND r.destination_airport_id = f.destination_airport_id
        WHERE r.route_id IS NULL
            AND f.airline_id IS NOT NULL
            AND f.origin_airport_id IS NOT NULL
            AND f.destination_airport_id IS NOT NULL;
        """
    )

    with ENGINE.begin() as con:
        result = con.execute(insert_sql)
        print(f" → Inserted {result.rowcount or 0} route rows")

    if compute_distance:
        compute_route_distances()

def compute_route_distances() -> None:
    """
    Compute approximate great-circle distance in nautical miles for each route.

    Uses airline.airports.latitude / longitude (degrees).
    """

```

```

print("◆ Computing distance_nm for airline.routes ...")

update_sql = text(
"""
UPDATE airline.routes r
SET distance_nm = sub.distance_nm::integer
FROM (
    SELECT
        r2.route_id,
        (
            2 * 6371 * asin(
                sqrt(
                    sin(radians(ad.latitude - ao.latitude) / 2)^2 +
                    cos(radians(ao.latitude)) *
                    cos(radians(ad.latitude)) *
                    sin(radians(ad.longitude - ao.longitude) / 2)^2
                )
            ) / 1.852
        ) AS distance_nm
    FROM airline.routes r2
    JOIN airline.airports ao
        ON ao.airport_id = r2.origin_airport_id
    JOIN airline.airports ad
        ON ad.airport_id = r2.destination_airport_id
    WHERE r2.distance_nm IS NULL
        AND ao.latitude IS NOT NULL
        AND ao.longitude IS NOT NULL
        AND ad.latitude IS NOT NULL
        AND ad.longitude IS NOT NULL
    ) AS sub
WHERE r.route_id = sub.route_id;
"""

)

with ENGINE.begin() as con:
    result = con.execute(update_sql)
    print(f" → Updated distance_nm for {result.rowcount or 0} routes")

def backfill_route_ids_on_flights() -> None:
    """Update airline.flights.route_id to match airline.routes."""

    print("◆ Backfilling flights.route_id from routes ...")

    update_sql = text(
"""
UPDATE airline.flights f
SET route_id = r.route_id
FROM airline.routes r
WHERE f.route_id IS NULL
    AND f.airline_id      = r.airline_id
    AND f.origin_airport_id = r.origin_airport_id
    AND f.destination_airport_id = r.destination_airport_id;
"""

)

    with ENGINE.begin() as con:
        result = con.execute(update_sql)
        print(f" → Updated route_id on {result.rowcount or 0} flights")

```

2. AIRCRAFT + FLIGHT.AIRCRAFT_ID

```

# -----
def backfill_aircraft() -> None:
    """
    Seed airline.aircraft with a small synthetic global fleet.

    Schema:
        aircraft_id      bigserial PK
        manufacturer    text
        model            text
        seat_capacity    integer
        tail_number     text (nullable)
    """

    print("◆ Backfilling airline.aircraft with synthetic fleet ...")

    insert_sql = text(
        """
        -- Only seed if table is currently empty
        INSERT INTO airline.aircraft (manufacturer, model, seat_capacity, tail_number)
        SELECT manufacturer, model, seat_capacity, tail_number
        FROM (
            VALUES
                ('Airbus', 'A320',      150, NULL),
                ('Boeing', '737-800',   165, NULL),
                ('Airbus', 'A321',      185, NULL),
                ('Boeing', '787-8',     242, NULL),
                ('Airbus', 'A350-900',   300, NULL)
            ) AS v(manufacturer, model, seat_capacity, tail_number)
        WHERE NOT EXISTS (SELECT 1 FROM airline.aircraft);
    """
    )

    with ENGINE.begin() as con:
        result = con.execute(insert_sql)
        print(f" → Inserted {result.rowcount or 0} aircraft rows")

def assign_aircraft_to_flights() -> None:
    """
    Assign an aircraft_id to each flight, picking randomly from all aircraft.
    (aircraft is not tied to a specific airline in your schema.)
    """

    print("◆ Assigning aircraft_id to flights ...")

    update_sql = text(
        """
        WITH choices AS (
            SELECT
                f.flight_id,
                (
                    SELECT ac2.aircraft_id
                    FROM airline.aircraft ac2
                    ORDER BY random()
                    LIMIT 1
                ) AS aircraft_id
            FROM airline.flights f
            WHERE f.aircraft_id IS NULL
        )
        UPDATE airline.flights f
        SET aircraft_id = c.aircraft_id
        FROM choices c
        WHERE f.flight_id = c.flight_id;
    """
)

```

```

    )

with ENGINE.begin() as con:
    result = con.execute(update_sql)
    print(f" → Updated aircraft_id on {result.rowcount or 0} flights")

# -----
# 3. FLIGHT_CHANGES
# -----


def generate_flight_changes(change_fraction: float = 0.05) -> None:
    """
    Create synthetic aircraft change events for a subset of flights.

    change_fraction: approximate fraction of flights to get a change row.
    """

    print(f"◆ Generating flight_changes for ~{change_fraction*100:.1f}% of flights ...")

    insert_sql = text(
        """
        WITH candidate_flights AS (
            SELECT f.flight_id, f.aircraft_id, f.flight_date
            FROM airline.flights f
            WHERE f.aircraft_id IS NOT NULL
        ),
        sampled AS (
            SELECT *
            FROM candidate_flights
            WHERE random() < {change_fraction}
        ),
        new_aircraft AS (
            SELECT
                s.flight_id,
                s.aircraft_id AS original_aircraft_id,
                s.flight_date,
                (
                    SELECT ac2.aircraft_id
                    FROM airline.aircraft ac2
                    WHERE ac2.aircraft_id <> s.aircraft_id
                    ORDER BY random()
                    LIMIT 1
                ) AS new_aircraft_id
            FROM sampled s
        )
        INSERT INTO airline.flight_changes (
            flight_id,
            old_aircraft_id,
            new_aircraft_id,
            reason,
            changed_at
        )
        SELECT
            n.flight_id,
            n.original_aircraft_id,
            n.new_aircraft_id,
            (ARRAY['Maintenance', 'Crew Reassignment', 'Operational', 'Equipment Downgrade'])[1
+ floor(random()*4)::text],
            n.flight_date::timestamp - INTERVAL '2 hours' + (random() * INTERVAL '90
minutes')
        FROM new_aircraft n
        WHERE n.new_aircraft_id IS NOT NULL;
        """
    )

```

```
)  
  
with ENGINE.begin() as con:  
    result = con.execute(insert_sql)  
    print(f" → Inserted {result.rowcount or 0} flight_changes rows")  
  
# Optionally: update flights.aircraft_id to reflect the "new" aircraft  
update_sql = text(  
    """  
        UPDATE airline.flights f  
        SET aircraft_id = fc.new_aircraft_id  
        FROM airline.flight_changes fc  
        WHERE fc.flight_id = f.flight_id;  
    """  
)  
  
with ENGINE.begin() as con:  
    result = con.execute(update_sql)  
    print(f" → Updated aircraft_id on {result.rowcount or 0} flights to match  
flight_changes")  
  
# -----  
# MAIN  
# -----  
  
def run():  
    print("== Backfill derived tables from flights ==")  
    backfill_routes_from_flights(compute_distance=False) # set True if you want distances  
computed  
    backfill_route_ids_on_flights()  
    backfill_aircraft()  
    assign_aircraft_to_flights()  
    generate_flight_changes(change_fraction=0.05)  
    print("== Done. ✅ ==")  
  
if __name__ == "__main__":  
    run()
```